

GAN Based Texture Rendering

Joon Jung (joonjung@stanford.edu)

Stanford University

Motivation

Rendering a physically realistic 3D model requires intensive computations, performing high resolution texture sampling and shading repeatedly. Among the intensive computations, the memory bandwidth, busy looping to read and write the textures, cost the most significant amount of resources.

Recently, various researches on extracting latent features from images and reconstructing the original from it have indeed progressed significantly. Now we can even create non-existing images based on the learned latent space features. [1] What if we can extract the low dimensional latent features of a high resolution 3D model, such as the light reflection distribution [3], and run through a GAN based generator to reconstruct the final 3D image? If so we might be able to completely skip the repetitive shading operations and avoid time consuming memory operations. We even might be able to just feed in the raw polygons and paint the textures without the 3D pipeline's help. Just maybe someday soon.

CyclGan

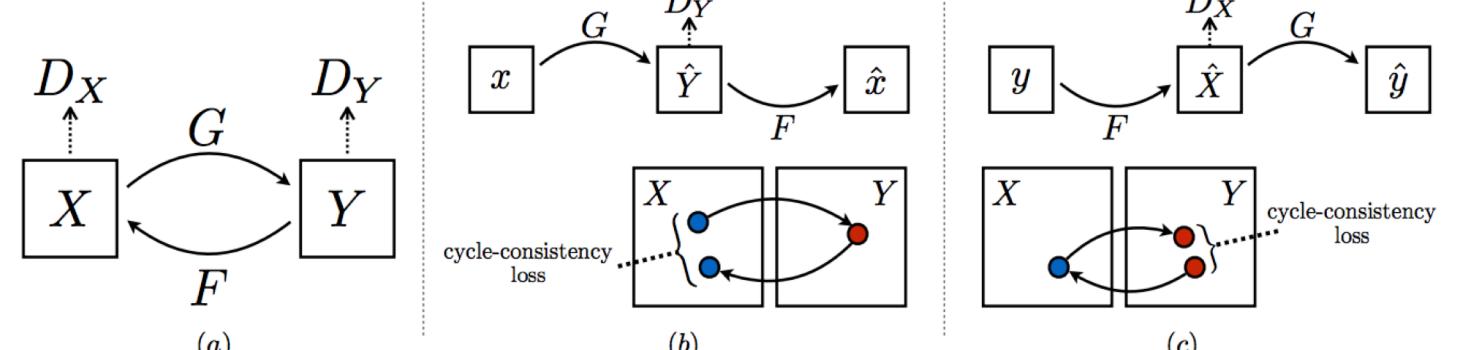


Figure 3: (a) Our model contains two mapping functions $G : X \rightarrow Y$ and $F : Y \rightarrow X$, and associated adversarial discriminators D_Y and D_X . D_Y encourages G to translate X into outputs indistinguishable from domain Y , and vice versa for D_X and F . To further regularize the mappings, we introduce two *cycle consistency losses* that capture the intuition that if we translate from one domain to the other and back again we should arrive at where we started: (b) forward cycle-consistency loss: $x \rightarrow G(x) \rightarrow F(G(x)) \approx x$, and (c) backward cycle-consistency loss: $y \rightarrow F(y) \rightarrow G(F(y)) \approx y$

CycleGan Model(excerpt from [1])

Losses

Adversarial Loss:

$$L_{GAN_G}(G, D_Y, X, Y) = \mathbb{E}_{x \sim p_{data}(x)}[(D_Y(G(x)) - 1)^2]$$

$$L_{GAN_D}(G, D_Y, X, Y) = \mathbb{E}_{y \sim p_{data}(y)}[(D_Y(y) - 1)^2] + \mathbb{E}_{x \sim p_{data}(x)}[(D_Y(G(x)))^2]$$

$$L_{cyc}(G, F) = \mathbb{E}_{x \sim p_{data}(x)}[||F(G(x)) - x||_1] + \mathbb{E}_{y \sim p_{data}(y)}[||G(F(y)) - y||_1]$$

Identity Loss:

$$L_{identity}(G, F) = \mathbb{E}_{x \sim p_{data}(x)}[||F(x) - x||_1] + \mathbb{E}_{y \sim p_{data}(y)}[||G(y) - y||_1]$$

Physically Based Rendering(PBR)

Reflectance Equation using Bidirectional Reflective Distribution Function (BRDF)

$$L_o(p, \omega_o) = \int f_r(p, \omega_i, \omega_o) L_i(p, \omega_i) n \cdot \omega_i d\omega_i$$

Incoming Radiance: $L_i(p, \omega_i)$

Incoming & Outgoing Solid Angle: ω_i, ω_o

Surface Normal: n

Hemisphere of Reflection: Ω

$$\text{BRDF Scale Factor: } f_r(p, \omega_i, \omega_o) = k_d \frac{c}{\pi} + \frac{DFG}{4(\omega_o \cdot n)(\omega_i \cdot n)}$$

D: MicroFacet Normal Distribution Function

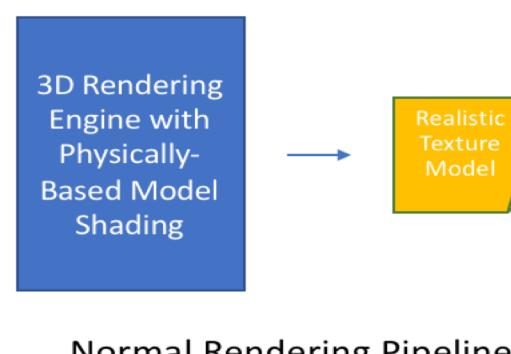
F: Fresnel Equation

G: Geometry Equation

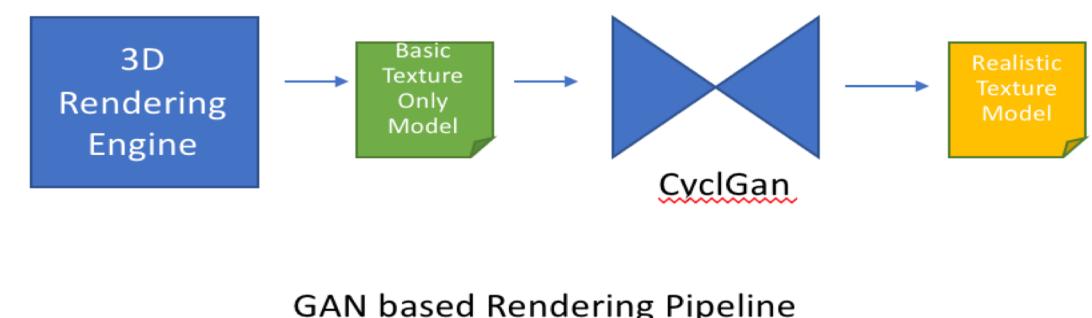


Left: an input image. Right: a target output image.

Model

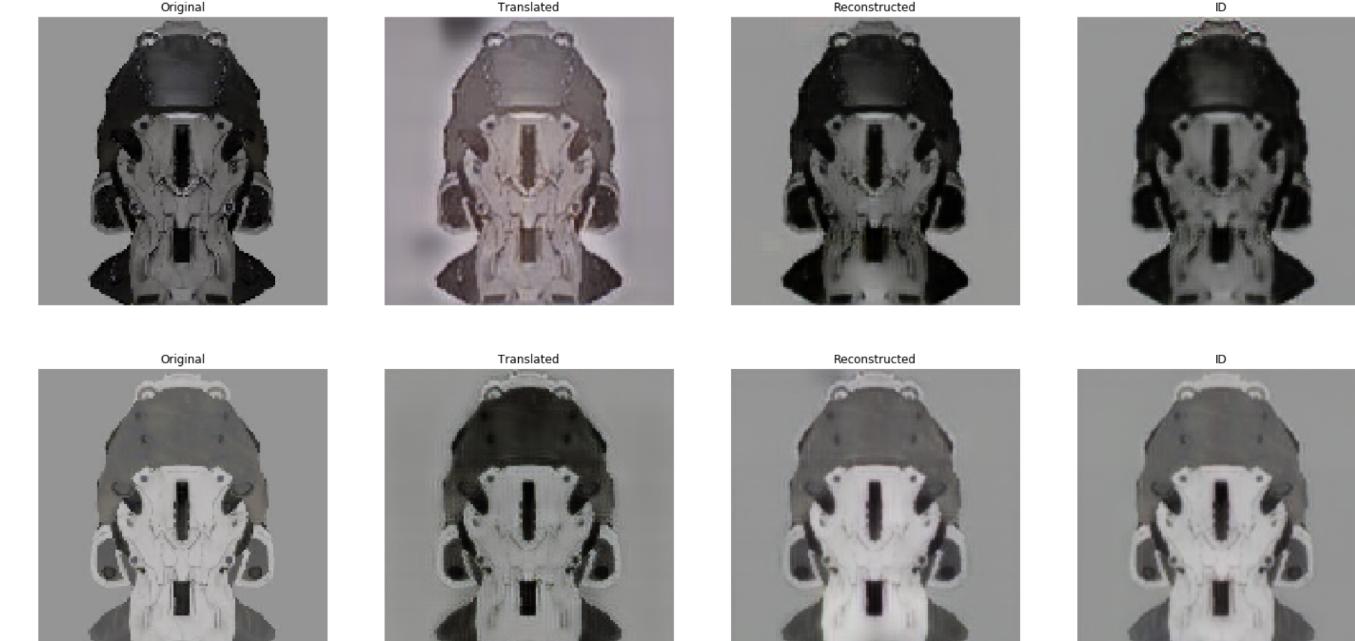


Normal Rendering Pipeline



GAN based Rendering Pipeline

Result Example



Losses

Discriminator Loss: 0.000191. Generator Loss: 2.440339

Error Analysis & Next Steps

- Mean Squared Error:** For the particular example shown above, the MSE for the pixel differences between the PBR rendered image and the Generator rendered image is **0.05527**. As a comparison, the MSE between the PBR and the reconstructed images is **0.00933**, almost 50 fold of differences.
- Human Eye Inspection:** Just by inspecting the generated image of F (row 2, 'Translated' on the figure above), we can see the light reflection is completely missing. This can be confirmed in the G generated image as well. In the generated image of G, you can see that it left out the light reflection intact for the generated image. This is not supposed to be so since the image represents non PBR, light reflection-less image.
- Next Steps:** For the remaining of the project, we will try to train the deep model to recognize the light reflection. We will separately train the generators by different light distributions (BRDF Scale Factor Function) and try to produce a better target output.

[1] Jun-Yan Zhu*, Taesung Park*, Phillip Isola, and Alexei A. Efros. "Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks", in IEEE International Conference on Computer Vision (ICCV), 2017. (* indicates equal contributions)

[2] Generative Deep Learning by David Foster. Publisher: O'Reilly Media, Inc. Release Date: June 2019 ISBN: 9781492041931 https://github.com/davidADSP/GDL_code

[3] Details on PBR: <https://learnopengl.com/PBR/Theory>

[4] Rendering Physically-Based ModelLO Materials August 5, 2018 by Warren Moore

References

Presentation URL

- Google Drive Link: https://drive.google.com/file/d/1nsnVWznsaqx4W_uiMyRRTETfjHACqg27/view?usp=sharing